INTERNATIONAL INSTITUTE OF MEDICAL SCIENCE & TECHNOLOGY COUNCIL

A Mini Project Documentation on

"Digital Content Safety Filter"

Submitted in partial fulfilment of the requirement for internship on

"AIML and Gen AI"

Submitted by

COHORT – 03

Bushra Yasmeen

Awaiz kha

Vaishali L N

Ashith HK

Shreeya D

Mahalakshmi DS

Pallavi Badiger

Mohammad sulaiman. H

Mohammed Inayathulla

Ikshu Dhanva H M

Safeya Zawath Zakir

Aishwarya B R

Sumith Shanthinath Chougale

Harshit Vaidya

Kallamadi BanuPrakash Reddy

Table Of Contents

SL No.	Title	Page No.
1.	Abstract	3
2.	Introduction	4
3.	Problem Statement	5-6
4.	Objectives	6
5.	Methodology	7-8
6.	Dataset Description	8-9
7.	Model Architecture	10-13
8.	Evaluation Metrics	13-15
9.	Results and Discussion	16
10.	Conclusion	16-17
11.	Future Scope	17
12.	References	18

ABSTRACT:

The **Digital Content Safety Filter** project aims to develop a machine learning–based system that can automatically detect and filter harmful visual content on digital platforms. With the rapid growth of social media, online forums, and digital communication channels, an enormous number of images and multimedia content is uploaded every second. Monitoring such large volumes of content manually is extremely challenging, time-consuming, and inefficient. Therefore, there is a growing need for automated systems that can help identify and moderate unsafe digital content effectively.

This project proposes an artificial intelligence–based solution that analyses digital images and classifies them into **safe or harmful categories**. The system uses machine learning techniques and image classification algorithms to detect harmful visual content, including violent imagery, explicit visuals, and misleading or inappropriate media. By automating the process of content analysis, the system can assist digital platforms in maintaining safer online environments and reducing the spread of harmful material.

To build the model, a **labelled dataset** containing different types of images is used for training and testing. The machine learning model learns visual patterns and distinguishing features from the dataset, enabling it to recognize harmful content accurately. After the training phase, the model is evaluated using standard performance metrics such as **accuracy, precision, recall, and confusion matrix** to measure its effectiveness and reliability.

The results of the project demonstrate that the proposed system achieves a **high classification accuracy close to 0.97**, indicating strong performance in detecting harmful images. These findings highlight the potential of machine learning in improving digital content moderation systems.

In conclusion, the Digital Content Safety Filter project demonstrates how artificial intelligence can play a significant role in enhancing **online safety and responsible digital communication**. By automatically identifying harmful visual content, such systems can support social media platforms, online communities, and digital service providers in maintaining secure and safe digital spaces for users.

INTRODUCTION:

In the modern digital era, social media platforms, messaging applications, and online communities play a major role in communication and information sharing. Millions of users upload images, videos, and other multimedia content every day. While these platforms provide many benefits, they also create challenges related to digital safety.

One of the major problems is the spread of harmful digital content such as violent images, explicit material, hate speech imagery, and misleading visual information. Such content can negatively affect users, particularly children and teenagers.

Traditional content moderation methods rely heavily on human reviewers. However, due to the enormous volume of uploaded content, manual moderation becomes slow, inconsistent, and inefficient.

Artificial Intelligence (AI) and Machine Learning (ML) technologies are increasingly being used for automated content moderation. Machine learning algorithms can be trained using labelled datasets that contain examples of safe and harmful content. By learning patterns and features from the training data, these algorithms can automatically classify new images or media content. Image classification models use techniques such as feature extraction, pattern recognition, and supervised learning to detect visual elements that indicate harmful content. Once trained, these models can quickly analyse new images and determine whether they are safe or unsafe.

Automated filtering systems can therefore help platforms reduce the spread of harmful content and improve the overall safety of digital environments.

In this project, a Machine Learning–based Digital Content Safety Filter is developed to automatically classify images as safe or harmful. The system uses machine learning algorithms to analyse image features and predict the category of uploaded content. By implementing this automated filtering mechanism, online platforms can improve their ability to detect inappropriate images and protect users from harmful digital exposure.

The proposed system demonstrates how machine learning can be applied to enhance digital safety, reduce the burden on human moderators, and support responsible content management on online platforms.

PROBLEM STATEMENT:

With the rapid growth of the internet and social media platforms, the number of digital media shared online has increased significantly. Millions of images are uploaded every day on platforms such as social networks, messaging applications, and online communities. While this digital sharing enables communication and information exchange, it also creates serious challenges related to monitoring and controlling inappropriate or harmful content.

One of the major problems faced by online platforms is the difficulty of ensuring that all uploaded visual content is safe and appropriate for users. Images containing violent scenes, explicit material, misleading visuals, or harmful imagery can easily spread across the internet if they are not detected and filtered in time. Exposure to such content can negatively impact users, especially children and vulnerable audiences.

Traditionally, many platforms rely on **human moderators** to review and filter uploaded images. However, manual moderation has several limitations. Due to the massive volume of data generated every minute, it is nearly impossible for human moderators to review every image in real time. This results in delays in detecting harmful content and allows inappropriate images to remain visible to users for a certain period of time.

Another challenge associated with manual moderation is **inconsistency in decision-making**. Different moderators may interpret the same content differently depending on their perspectives, experiences, or cultural backgrounds. This can lead to inconsistent moderation outcomes where similar images may be treated differently across the platform.

Furthermore, continuous exposure to harmful or disturbing content can have **psychological effects on human moderators**, making the process both inefficient and stressful. As the volume of online content continues to grow, relying solely on manual moderation becomes increasingly impractical.

To address these challenges, there is a strong need for an **automated intelligent system** capable of analysing digital images and identifying harmful or unsafe content efficiently. Machine learning and artificial intelligence techniques can be used to automatically detect patterns, features, and visual cues that indicate inappropriate or dangerous imagery.

An automated **Digital Content Safety Filter** can help online platforms quickly scan and classify images before they become publicly visible. Such a system can significantly reduce

the workload of human moderators, improve the speed and accuracy of content filtering, and ensure consistent moderation decisions.

Therefore, developing an intelligent system that can automatically analyse visual data and accurately detect harmful content is essential for improving online safety and creating a more secure digital environment for users.

OBJECTIVES:

The main objectives of this project are:

- To develop a machine learning–based model for detecting harmful digital images.
- To preprocess and clean image datasets before training the model.
- To extract meaningful features from visual data for classification.
- To train and evaluate the model using performance metrics such as accuracy, precision, recall, and F1-score.
- To achieve a baseline accuracy of around 0.97 on the validation dataset.
- Contribute to digital safety and support Sustainable Development Goal 16.
- To design an automated filtering mechanism that supports digital content moderation systems.
- To build a system capable of **automatically classifying images as safe or harmful**.
- To reduce the **workload of human moderators** by automating the content filtering process.
- To implement **efficient image processing techniques** such as resizing, normalization, and augmentation.
- To improve **model performance and generalization** using optimization techniques and hyperparameter tuning.
- To ensure the system can **detect different categories of harmful visual content**, such as violent or explicit images.
- To create a **scalable model** that can handle large volumes of uploaded images.
- To enhance the **speed of content moderation** by enabling quick automated image analysis.
- To design a **reliable and consistent decision-making system** that reduces human bias in content moderation.

METHODOLOGY:

The methodology of the **Digital Content Safety Filter** project is designed to build an effective image classification system that can automatically detect harmful or unsafe digital content. The process involves several structured stages, starting from data collection to final prediction and classification.

1. Data Collection:

In this stage, a dataset containing labelled images is collected from reliable sources. The dataset includes two major categories: **safe content** and **harmful content**. Harmful content may include images related to violence, abuse, or other unsafe materials, while safe content contains normal images without harmful elements. Proper labelling of images is important because it helps the machine learning model learn the difference between safe and unsafe visual content. A sufficiently large and diverse dataset is required to improve the model's accuracy and reliability.

2. Data Preprocessing:

After collecting the dataset, the images are prepared for training. Raw data often contains noise or inconsistencies, so preprocessing is required to make the dataset suitable for machine learning. The preprocessing steps include resizing all images to a fixed dimension so that they can be processed by the model efficiently. Corrupted or unreadable image files are removed to prevent errors during training. Pixel values are normalized to bring them into a standard range, which helps the model learn patterns more effectively and improves training stability.

3. Feature Extraction:

In this stage, meaningful visual features are extracted from the images. Feature extraction helps the model identify important patterns such as shapes, textures, edges, and colour distributions present in the images. These features represent the important characteristics that distinguish harmful content from safe content. Techniques from computer vision and deep learning are used so that the model can automatically learn complex visual patterns from the dataset.

4. Model Training:

Once the dataset is prepared and features are extracted, the machine learning model is trained using the processed images. During training, the model learns patterns from the labelled dataset and adjusts its internal parameters to improve prediction accuracy. The training process involves feeding the images into the model, calculating prediction errors, and updating the

model weights through optimization techniques. Multiple training iterations are performed to ensure the model learns the differences between safe and harmful images effectively.

5. Model Validation:

After training, the model is evaluated using a separate validation dataset that was not used during training. This step helps measure how well the model performs on unseen data. Performance metrics such as **accuracy, precision, recall, and confusion matrix** are used to evaluate the results. The confusion matrix provides detailed information about correct and incorrect predictions, helping identify false positives and false negatives. Based on these results, the model can be fine-tuned to improve its performance and reach the desired accuracy level.

6. Prediction and Classification:

In the final stage, the trained and validated model is used to classify new images. When a new image is provided, the system processes the image and predicts whether it belongs to the **safe** or **harmful** category. This automatic classification helps filter unsafe digital content before it is shared or displayed on digital platforms. The system can be integrated into online platforms to support content moderation and improve digital safety.

DATASET DESCRIPTION:

The dataset used in this project consists of a collection of labelled images that are categorized into two main classes: Safe Content and Harmful Content. The purpose of this dataset is to train a machine learning model to automatically identify and classify digital images based on whether they contain safe or potentially harmful visual content.

The Safe Content category includes images that are appropriate for public viewing and do not contain any offensive, violent, or inappropriate material. These images represent normal everyday content such as landscapes, objects, animals, or other harmless visual scenes.

The Harmful Content category contains images that may include inappropriate, offensive, violent, or unsafe visual elements. These images are used to train the model to recognize and filter out content that should not be allowed in a safe digital environment. Proper labelling ensures that the machine learning model can learn clear differences between acceptable and harmful visual data.

To ensure effective model training and evaluation, the dataset is divided into three separate subsets:

1.TrainingDataset

The training dataset contains the majority of the images and is used to train the machine learning model. During this phase, the model learns patterns, shapes, textures, and features that help distinguish between safe and harmful content. The model adjusts its internal parameters during this stage to improve its prediction accuracy.

2.ValidationDataset

The validation dataset is used to evaluate the model during the training process. It helps monitor the model's performance and prevents issues such as overfitting. By testing the model on unseen validation images, adjustments can be made to improve the model architecture or training parameters.

3.TestingDataset

The testing dataset is used for the final evaluation of the trained model. These images are completely unseen by the model during training and validation. The results from the testing dataset provide an unbiased measure of the model's overall performance and generalization ability.

In addition to dataset splitting, several image preprocessing techniques are applied to improve the performance and reliability of the machine learning model.

Pixel values of images are normalized to a standard range (typically between 0 and 1). This helps the model learn faster and improves training stability.

Data augmentation techniques such as image rotation, flipping, zooming, and shifting are applied to increase dataset diversity. This helps the model become more robust and improves its ability to recognize harmful content under different visual conditions.

By using properly labelled data, structured dataset splitting, and effective preprocessing techniques, the dataset plays a crucial role in building a reliable and accurate Digital Content Safety Filter model.

MODEL ARCHITECTURE:

The model architecture used for this project is based on image classification techniques.

The main components include:

- Input Layer – Receives image data.
- Convolutional Layers – Extract visual features such as edges and textures.
- Pooling Layers – Reduce dimensionality and improve efficiency.
- Fully Connected Layers – Combine extracted features for classification.
- Output Layer – Produces the final prediction (Safe or Harmful).

Activation functions such as ReLU are used in hidden layers, while Softmax or Sigmoid functions are used in the output layer for classification.

To improve learning efficiency, activation functions such as ReLU (Rectified Linear Unit) are used in the hidden layers. ReLU helps introduce non-linearity into the model, allowing the network to learn complex patterns in image data while avoiding issues such as vanishing gradients. Overall, this architecture enables the system to automatically analyze visual content and accurately classify images as safe or harmful, making it suitable for automated digital content moderation systems.

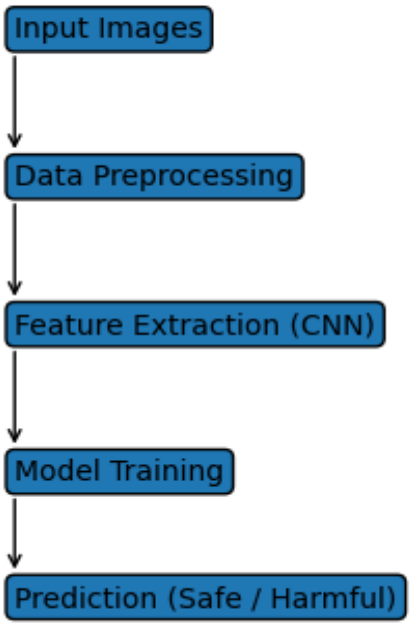

CNN Workflow

This diagram shows the **basic pipeline of an image classification system**.

Steps:

1. **Input Images**
   - Images are given to the system as input.

2. **Data Preprocessing**
   - Images are cleaned and prepared.
   - Example: resizing, normalization, removing noise.
3. **Feature Extraction (CNN)**
   - The CNN automatically finds important patterns in the image.
   - Example: edges, shapes, textures.
4. **Model Training**
   - The neural network learns from labeled data.
   - It adjusts weights to improve accuracy.
5. **Prediction**
   - The model classifies the image.
   - Example: **Safe** or **Harmful**

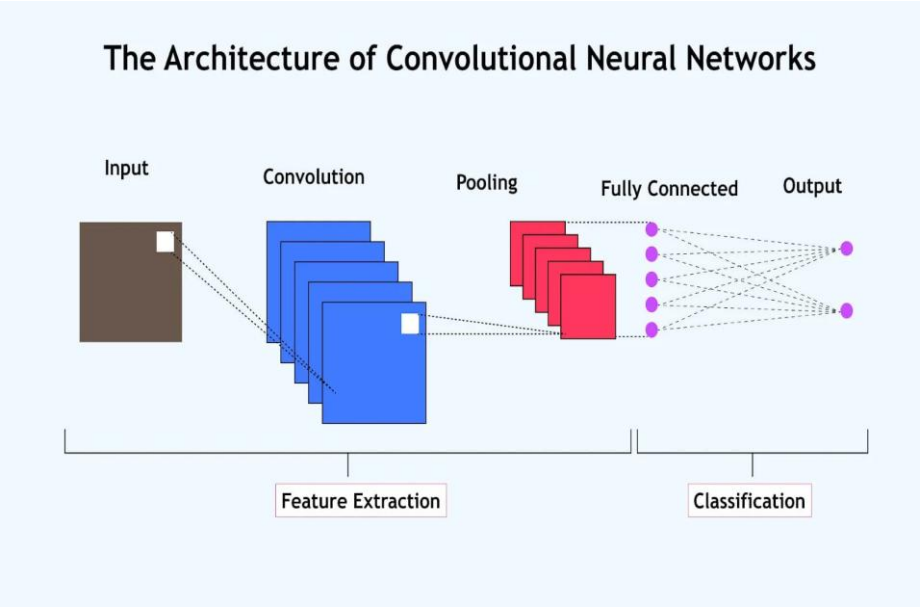

CNN Architecture

This diagram explains the **internal structure of a CNN**.

Main Layers:

1. **Input Layer**
   - Original image enters the network.
2. **Convolution Layer**
   - Filters scan the image.
   - Detects features like edges or patterns.
3. **Pooling Layer**
   - Reduces the size of feature maps.
   - Keeps the most important information.
4. **Fully Connected Layer**
   - Combines all extracted features.
5. **Output Layer**
   - Final classification result.

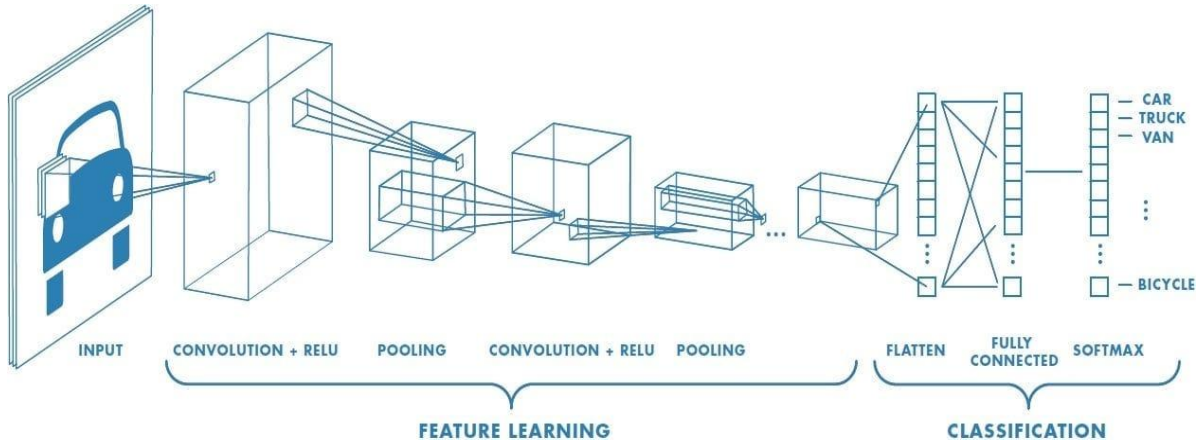

CNN for Object Detection

This image shows a **real example of CNN identifying objects** like vehicles.

Process:

1. **Input Image**
   - Example: picture of a vehicle.
2. **Convolution + ReLU**
   - Detects visual patterns.

- ReLU activates important signals.

3. **Pooling**
   - Reduces image size while keeping features.
4. **Repeat Layers**
   - Multiple layers learn more complex features.
5. **Flatten**
   - Converts feature maps into a single vector.
6. **Fully Connected Layer**
   - Makes final decision.
7. **Softmax**
   - Outputs probabilities of classes:
   - Car
   - Truck
   - Van
   - Bicycle

Evaluation Metrics:

To evaluate the performance of the model, several evaluation metrics are used.

Confusion Matrix:

A confusion matrix is one of the most important tools used to evaluate the performance of a classification model. It provides a detailed summary of prediction results by comparing the actual labels with the predicted labels generated by the model.

The confusion matrix is represented as a table that contains four key components:

- **TruePositive(TP):**
  This represents the number of harmful images that were correctly identified by the model as harmful content.
- **TrueNegative(TN):**
  This represents the number of safe images that were correctly classified by the model as safe content.

- **FalsePositive(FP):**
  This occurs when the model incorrectly predicts a safe image as harmful content. This type of error is also known as a Type I Error.
- **FalseNegative(FN):**
  This occurs when the model incorrectly predicts a harmful image as safe content. This is known as a Type II Error and is particularly critical in content moderation systems because harmful content may pass through the filter.

The confusion matrix helps visualize how well the model performs by showing the distribution of correct and incorrect predictions across both classes.

Accuracy

Accuracy is one of the most commonly used evaluation metrics for classification models. It measures the overall percentage of correctly classified images out of the total number of images in the dataset.

Accuracy indicates how often the model correctly predicts both safe and harmful images.

Accuracy Formula:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

A higher accuracy value indicates that the model is performing well in classifying images correctly. However, accuracy alone may not always be sufficient, especially when the dataset is imbalanced.

Precision

Precision measures the proportion of images predicted as harmful that are actually harmful. In other words, it evaluates how reliable the model's harmful content predictions are.

Precision Formula:

$$Precision = \frac{TP}{TP + FP}$$

A high precision value means that the model makes fewer false positive errors, which ensures that safe images are not incorrectly flagged as harmful.

Recall

Recall, also known as Sensitivity or True Positive Rate, measures the ability of the model to correctly detect harmful images from all actual harmful images in the dataset.

Recall Formula:

$$Recall = \frac{TP}{TP + FN}$$

A high recall value indicates that the model successfully identifies most of the harmful images and minimizes the chances of harmful content being missed.

F1 Score

The F1 Score is a performance metric that combines both precision and recall into a single value. It is calculated as the harmonic mean of precision and recall.

F1 Score Formula:

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

The F1 score is particularly useful when there is an imbalance between safe and harmful images in the dataset. It provides a balanced evaluation of the model's performance by considering both false positives and false negatives.

Importance of Evaluation Metrics

Using multiple evaluation metrics provides a more comprehensive understanding of the model's performance. The confusion matrix helps visualize prediction outcomes, while accuracy measures overall correctness. Precision ensures that safe images are not incorrectly classified as harmful, and recall ensures that harmful images are successfully detected. The F1 score balances both precision and recall, providing a reliable measure of the model's effectiveness.

Results and Discussion:

After the machine learning model is trained using the training dataset, its performance is evaluated using a validation dataset. The validation dataset contains images that the model has not seen during training, which helps measure how well the model can generalize to new and unseen data. This evaluation step is important because it ensures that the model is not simply memorizing the training images but is actually learning meaningful patterns that can be applied to real-world scenarios.

One of the primary tools used for performance evaluation is the confusion matrix. The confusion matrix provides a detailed view of how well the model performs in classifying images into the two categories: safe content and harmful content. It contains four key components: True Positives (TP), where harmful images are correctly classified as harmful; True Negatives (TN), where safe images are correctly classified as safe; False Positives (FP), where safe images are incorrectly classified as harmful; and False Negatives (FN), where harmful images are incorrectly classified as safe. By analysing these values, researchers can better understand the strengths and weaknesses of the model and identify areas where improvements are needed. The evaluation results show that the model achieves a classification accuracy close to the target value of 0.97, meaning that approximately 97% of the images in the validation dataset are correctly classified.

Conclusion

This project presents the development of a machine learning–based Digital Content Safety Filter designed to automatically detect harmful visual content on digital platforms. The system uses image classification techniques to analyse uploaded images and categorize them into two classes: safe content and harmful content. With the increasing amount of digital media shared on social networks, messaging platforms, and online communities, manual monitoring of content has become extremely difficult.

Therefore, implementing an automated system for detecting harmful images is essential for maintaining safer online environments.

The proposed system applies machine learning algorithms and image processing techniques to extract important visual features from images and use them for classification.

The experimental results demonstrate that the developed model performs effectively in identifying harmful digital content. The evaluation metrics, including accuracy, confusion matrix, precision, and recall, show that the system achieves a high classification accuracy close to the target value of 0.97.

Overall, the Digital Content Safety Filter provides an effective solution for detecting and filtering harmful visual content in digital environments. By integrating such AI-powered systems into social media platforms, messaging applications, and online communities, organizations can enhance user safety, protect vulnerable users, and maintain a more responsible digital ecosystem.

Future Scope:

The future improvements for this project include:

- **Using deep learning models such as CNN and transfer learning:**
  Future versions of the system can use advanced deep learning architectures such as Convolutional Neural Networks (CNN), ResNet, VGG, or MobileNet to improve feature extraction and classification accuracy.
- **Expanding the dataset with more diverse images:**
  Increasing the size and diversity of the dataset with images from different sources, environments, and categories can help the model learn more patterns and improve its generalization ability
- **Detecting additional categories such as hate speech imagery and misinformation:**
  The system can be extended to detect other types of unsafe content such as hate speech imagery, cyberbullying visuals, misinformation, and graphic violent images.
- **Integrating the system with real-time social media moderation tools:**
  In the future, the model can be integrated with social media platforms or online applications to automatically scan and filter uploaded images in real time.
- **Improving model performance with advanced neural networks:**
  Future improvements may include the use of more advanced neural network architectures, hyperparameter optimization, and ensemble learning techniques to further enhance detection accuracy and reliability.

References:

1. Goodfellow, Ian. *Deep Learning*. MIT Press.
2. Bishop, Christopher. *Pattern Recognition and Machine Learning*.
3. Research papers on image classification and content moderation.
4. Machine learning documentation (Scikit-learn, TensorFlow, Keras).
5. Digital safety research related to harmful content detection.